

## BDA 551-Assignment 2

Feray Ece Topcu

July 14, 2018

- Load required libraries:

```
library(tree)
```

```
library(ISLR)
```

- Read Dataset & examine columns:

```
#attach(Carseats)
```

```
carseats <- Carseats
```

```
str(carseats)
```

```
## 'data.frame': 400 obs. of 11 variables:
```

```
## $ Sales : num 9.5 11.22 10.06 7.4 4.15 ...
```

```
## $ CompPrice : num 138 111 113 117 141 124 115 136 132 132 ...
```

```
## $ Income : num 73 48 35 100 64 113 105 81 110 113 ...
```

```
## $ Advertising: num 11 16 10 4 3 13 0 15 0 0 ...
```

```
## $ Population : num 276 260 269 466 340 501 45 425 108 131 ...
```

```
## $ Price : num 120 83 80 97 128 72 108 120 124 124 ...
```

```
## $ ShelfLoc : Factor w/ 3 levels "Bad","Good","Medium": 1 2 3 3 1 1 3 2  
3 3 ...
```

```
## $ Age : num 42 65 59 55 38 78 71 67 76 76 ...
```

```
## $ Education : num 17 10 12 14 13 16 15 10 10 17 ...
```

```
## $ Urban : Factor w/ 2 levels "No","Yes": 2 2 2 2 2 1 2 2 1 1 ...
```

```
## $ US : Factor w/ 2 levels "No","Yes": 2 2 2 2 1 2 1 2 1 2 ...
```

### Question 1. Split the data set into a training set and a test set.

```
set.seed(1)
train = sample(1:nrow(carseats), nrow(carseats)/4)
```

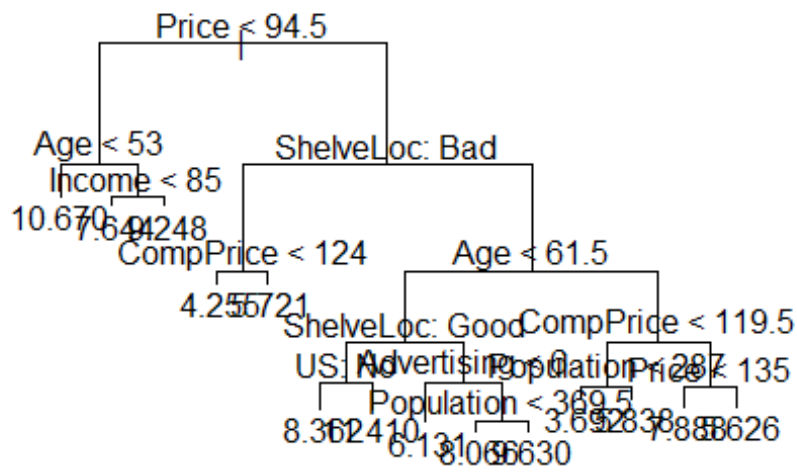
### Question 2. Fit a regression tree to the training set. Plot the tree, and interpret the results. What test error rate do you obtain ?

- Fit a tree with train dataset, plot the regression tree:

```
tree.carseats = tree(Sales~., carseats, subset = train) # fit tree
summary(tree.carseats)

##
## Regression tree:
## tree(formula = Sales ~ ., data = carseats, subset = train)
## Variables actually used in tree construction:
## [1] "Price"      "Age"      "Income"    "ShelveLoc" "CompPrice"
## [6] "US"        "Advertising" "Population"
## Number of terminal nodes: 14
## Residual mean deviance: 2.524 = 217.1 / 86
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -3.53200 -0.96410 -0.04806  0.00000  0.81960  4.41900

plot(tree.carseats)
text(tree.carseats, pretty = 0)
```




```

yhat = predict(tree.carseats, newdata = carseats[-train, ]) # test set predictions
table(yhat) # test set predictions

## yhat
##          3.692 4.254545454545455          5.626          5.72125
##          13          23          13          34
##          5.838 6.131111111111111 7.64428571428572 7.888333333333333
##           8          43          17          32
## 8.065555555555556          8.362          9.248          9.63
##          30           5           9          18
## 10.66777777777778          11.414
##          24          31

carseats.test = carseats[-train, "Sales"] # test set actual values
plot(yhat, carseats.test) # note as we average at each node predictions are b
unchd
abline(0, 1)

```



```
mean((yhat - carseats.test)^2) # calc MSE
```

```
## [1] 5.179786
```

- MSE is 5.179 when seed=1.

**Question 3. Use cross-validation in order to determine the optimal level of tree complexity. Does pruning the tree improve the test error rate ?**

- Apply cross validation:

```
cv.carseats=cv.tree(tree.carseats)
```

```
cv.carseats
```

```
## $size
```

```
## [1] 14 13 12 11 10 9 8 7 6 5 4 3 2 1
```


```
##
```

```
## $dev
```

```
## [1] 645.2752 671.5432 675.3820 688.7686 684.7958 686.0333 688.7464
```

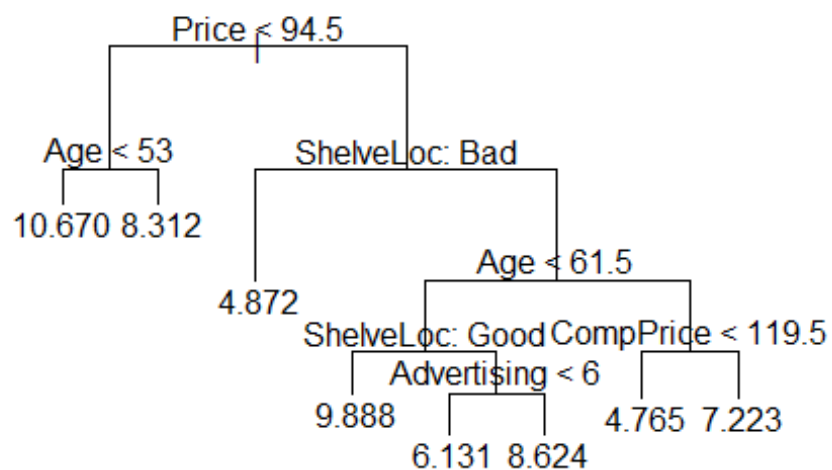
```
## [8] 721.4186 716.5108 706.1968 698.9934 635.1100 687.1329 706.3378
##
## $k
## [1] -Inf 7.501374 7.866921 9.963556 11.513290 18.064066
## [7] 23.286760 28.529143 34.052429 34.949434 38.038916 60.275200
## [13] 93.727701 104.867337
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"

plot(cv.carseats$size,cv.carseats$dev,type='b')
```




- Fitting Regression Trees using 8 terminal nodes:

```
prune.carseats=prune.tree(tree.carseats,best=8)
plot(prune.carseats)
text(prune.carseats,pretty=0)
```



- Calculate MSE for pruned tree:

```
yhat=predict(prune.carseats,newdata=carseats[-train,])
carseats.test=carseats[-train,"Sales"]
plot(yhat,carseats.test)
abline(0,1)
```



```
mean((yhat-carseats.test)^2)
```

```
## [1] 5.152601
```

MSE is 4.780099 for pruned tree. So, pruned tree gives better results.

**Question 4. Use the bagging approach in order to analyze this data. What test error rate do you obtain? Use the “importance()” function to determine which variables are most important. Use ntree=1000.**

- Apply bagging:

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1)
```

```

bag.carseats=randomForest(Sales~.,data=carseats,subset=train,mtry=13,ntree=10
00,importance=TRUE)

yhat.bag = predict(bag.carseats,newdata=carseats[-train,])
mean((yhat.bag-carseats.test)^2)

## [1] 3.57208

importance(bag.carseats)


##           %IncMSE IncNodePurity
## CompPrice  19.568050      97.703899
## Income      8.046965      43.960119
## Advertising 11.982097      62.940157
## Population  -1.348149      45.458754
## Price       44.934798     193.698574
## ShelfLoc    29.040872      94.287851
## Age         15.024811      82.350640
## Education   2.151313      25.327554
## Urban       2.355167       4.435798
## US          4.654712       9.741332

varImpPlot(bag.carseats)

```



### bag.carseats



MSE is 3.51809 for bagging. So, Bagging is a better algorithm for this dataset to forecast Sales. Due to output of importance() function, Price is the most important predictor in this dataset.

**Question 5. Use random forests to analyze this data. What test error rate do you obtain? Use the “importance()” function to determine which variables are most important. Use ntree=1000, mtry=3.**

```
library(randomForest)
set.seed(1)

bag.carseats=randomForest(Sales~.,data=carseats,subset=train,mtry=3,ntree=1000,importance=TRUE)
```

```
yhat.bag = predict(bag.carseats,newdata=carseats[-train,])
mean((yhat.bag-carseats.test)^2)
```


```
## [1] 4.417369
```

```
importance(bag.carseats)
```

##		%IncMSE	IncNodePurity
##	CompPrice	9.9937219	79.21195
##	Income	7.5513798	67.11762
##	Advertising	7.2126094	57.66589
##	Population	-3.4882346	57.20082
##	Price	29.6031283	139.44304
##	ShelveLoc	21.8266814	75.19197
##	Age	13.4348826	86.63961
##	Education	0.4021779	39.98516
##	Urban	1.5420490	10.29001
##	US	6.4859709	16.54111

```
varImpPlot(bag.carseats)
```

**bag.carseats**



MSE is 4.417369 for RandomForest while mtry=3 and ntree=1000. So, Bagging is a better method than RandomForest to forecast Sales.

**Question 6. By using 10 fold cross validation and grid search detect best parameters of ntree and mtry for random forests. What test error rate do you obtain by using best parameters.**

- GridSearch & CV:

```
library(caret)

## Loading required package: lattice

## Loading required package: ggplot2

##
## Attaching package: 'ggplot2'

## The following object is masked from 'package:randomForest':
##
##      margin


#library("FactoMineR")
#library("e1071")
#library(dplyr)

set.seed(1)

fitControl <- trainControl(## 10-fold CV
  method = "repeatedcv",
  number = 10,
  repeats = 1)
#

rf_gridsearch <- train(Sales~., data=carseats, method = "rf",
  trControl = fitControl, verbose = FALSE,
  tuneLength = 4)
```

```
plot(rf_gridsearch)
```



```
#plot(rf_gridsearch, plotType = "level")
```

```
rf_gridsearch$results
```

```
##      mtry      RMSE Rsquared      MAE      RMSESD RsquaredSD      MAESD
## 1      2 1.827595 0.6803037 1.474191 0.2203589 0.09071373 0.1609274
## 2      5 1.569929 0.7223584 1.260126 0.2174363 0.07556059 0.1707563
## 3      8 1.531404 0.7215676 1.221723 0.2129464 0.07542307 0.1723837
## 4     11 1.521658 0.7198204 1.218203 0.2038990 0.07049754 0.1700340
```

```
best(rf_gridsearch$results, metric="Rsquared", maximize=T)
```

```
## [1] 2
```

```
tolerance(rf_gridsearch$results, metric="Rsquared", maximize=T, tol=2)
```

```
## [1] 2
```

```
rf_gridsearch$results[2,]
```

##	mtry	RMSE	Rsquared	MAE	RMSESD	RsquaredSD	MAESD
## 2	5	1.569929	0.7223584	1.260126	0.2174363	0.07556059	0.1707563

The best value is 5 for mtry parameter. So, apply RandomForest with the best:

```
library(randomForest)
set.seed(1)

bag.carseats=randomForest(Sales~.,data=carseats,subset=train,mtry=5,ntree=1000,importance=TRUE)

yhat.bag = predict(bag.carseats,newdata=carseats[-train,])
mean((yhat.bag-carseats.test)^2)


## [1] 3.864939

importance(bag.carseats)

##           %IncMSE IncNodePurity
## CompPrice  17.515685      83.072362
## Income      6.715830      58.983759
## Advertising 9.722781      61.356736
## Population -1.734053      48.737245
## Price      35.898562     162.573253
## ShelfLoc   27.544971      87.295701
## Age       16.239920      86.265179
## Education   5.290071      34.353653
## Urban       2.095135       7.503634
## US         5.961277     13.292080

varImpPlot(bag.carseats)
```

### bag.carseats



MSE is 3.864939 when  $mtry=5$  and  $ntree=1000$  for RandomForest which is much more better before tuning. Before tuning, RandomForest's MSE is 4.4173 and MSE of bagging is 3.51809. Due to the results; tuned RandomForest is better than first RandomForest while Bagging approach is a bit improved than tuned RandomForest.